

Information Retrieval

Index Construction

Hamid Beigy

Sharif university of technology

October 6, 2018



Table of contents

1. Introduction
2. Sort-based index construction
3. Single-pass in-memory indexing (SPIMI)
4. Distributed indexing
5. Dynamic indexing



Table of contents

- 1 Introduction
- 2 Sort-based index construction
- 3 Single-pass in-memory indexing (SPIMI)
- 4 Distributed indexing
- 5 Dynamic indexing



Inverted index

- 1 The goal is constructing inverted index

For each term t , we store a list of all documents that contain t .

BRUTUS	→	1	2	4	11	31	45	173	174
--------	---	---	---	---	----	----	----	-----	-----

CAESAR	→	1	2	4	5	6	16	57	132	...
--------	---	---	---	---	---	---	----	----	-----	-----

CALPURNIA	→	2	31	54	101
-----------	---	---	----	----	-----

⋮

dictionary

postings



RCV1 collection

- 1 Shakespeare's collected works are not large enough for demonstrating many of the points in this course.
- 2 As an example for applying scalable index construction algorithms, we will use the Reuters **RCV1** collection.
- 3 English newswire articles sent over the wire in 1995 and 1996 (a year).
- 4 RCV1 statistics
 - Number of documents (N): **800,000**
 - Number of tokens per document (L): **200**
 - terms (M) : **400,000**
 - Bytes per token (including spaces): **6**
 - Bytes per token (without spaces): **4.5**
 - Bytes per term: **7.5**
- 5 Why does the algorithm given in previous sections **not scale** to **very large collections**?



Hardware Basics

- 1 Access to data is much **faster in memory than on disk**. (roughly a factor of 10)
- 2 **Disk seeks are "idle" time**: No data is transferred from disk while the disk head is being positioned.
- 3 To optimize transfer time from disk to memory: **one large chunk is faster than many small chunks**.
- 4 **Disk I/O is block-based**: Reading and writing of entire blocks (as opposed to smaller chunks). Block sizes: 8KB to 256 KB
- 5 Servers used in IR systems typically have **many GBs of main memory** and **TBs of disk space**.
- 6 **Fault tolerance is expensive**: Its cheaper to use many regular machines than one fault tolerant machine.



Hard Disk







Table of contents

- 1 Introduction
- 2 Sort-based index construction**
- 3 Single-pass in-memory indexing (SPIMI)
- 4 Distributed indexing
- 5 Dynamic indexing



Sort-based index construction

- 1 As we build index, we parse docs one at a time.
- 2 The final postings for any term are incomplete until the end.
- 3 Can we keep all postings in memory and then do the sort in-memory at the end?
- 4 **No**, not for large collections
- 5 Thus: We need to store intermediate results on disk.
- 6 Can we use the same index construction algorithm for larger collections, but by using disk instead of memory?
- 7 **No**: Sorting very large sets of records on disk is too slow– too many disk seeks.
- 8 We need an **external sorting** algorithm.

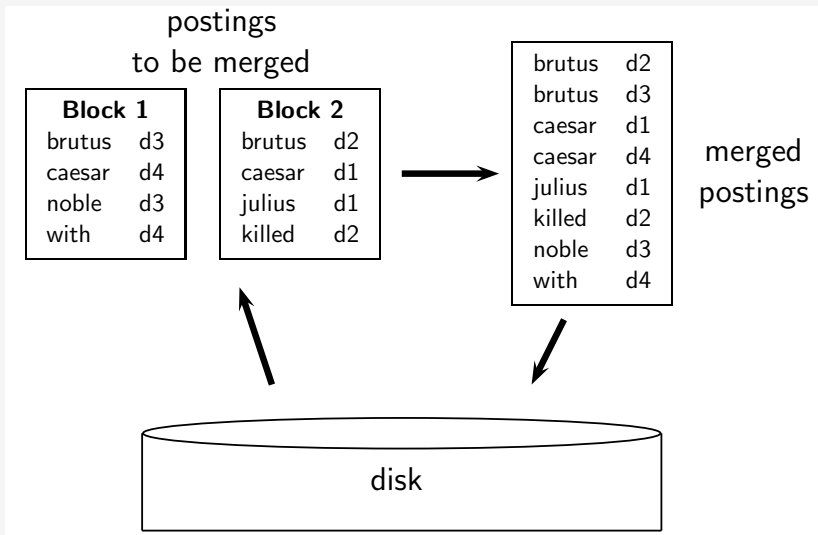


External sorting algorithm

- 1 We must sort $T = 100,000,000$ non-positional postings.
- 2 Each posting has size 12 bytes (4+4+4: termID, docID, term frequency).
- 3 Define a block to consist of $10,000,000$ such postings
- 4 We can easily fit that many postings into memory. We will have 10 such blocks for RCV1.
- 5 Basic idea of algorithm:
- 6 For each block do
 - accumulate postings
 - sort in memory
 - write to disk
- 7 Then merge the blocks into one long sorted order.



Merging two blocks





Blocked Sort-Based Indexing

BSBINDEXCONSTRUCTION()

```
1   $n \leftarrow 0$ 
2  while (all documents have not been processed)
3  do  $n \leftarrow n + 1$ 
4       $block \leftarrow \text{PARSENEXTBLOCK}()$ 
5       $\text{BSBI-INVERT}(block)$ 
6       $\text{WRITEBLOCKTODISK}(block, f_n)$ 
7   $\text{MERGEBLOCKS}(f_1, \dots, f_n; f_{\text{merged}})$ 
```



Problem with sort-based algorithm

- 1 The assumption was: **we can keep the dictionary in memory.**
- 2 We need the dictionary (which grows dynamically) in order to implement a term to termID mapping.
- 3 Actually, we could work with term,docID postings instead of termID,docID postings . . .
- 4 The **intermediate files become very large.** (We would end up with a scalable, but very slow index construction method.)



Table of contents

- 1 Introduction
- 2 Sort-based index construction
- 3 Single-pass in-memory indexing (SPIMI)**
- 4 Distributed indexing
- 5 Dynamic indexing



Single-pass in-memory indexing (SPIMI)

- 1 Key idea 1: Generate separate dictionaries for each block no need to maintain term-termID mapping across blocks.
- 2 Key idea 2: Dont sort. Accumulate postings in postings lists as they occur.
- 3 With these two ideas we can generate a complete inverted index for each block.
- 4 These separate indexes can then be merged into one big index.



Single-pass in-memory indexing algorithm

```
SPIMI-INVERT(token_stream)
1  output_file  $\leftarrow$  NEWFILE()
2  dictionary  $\leftarrow$  NEWHASH()
3  while (free memory available)
4  do token  $\leftarrow$  next(token_stream)
5      if term(token)  $\notin$  dictionary
6          then postings_list  $\leftarrow$  ADDTODICTIONARY(dictionary, term(token))
7          else postings_list  $\leftarrow$  GETPOSTINGSLIST(dictionary, term(token))
8      if full(postings_list)
9          then postings_list  $\leftarrow$  DOUBLEPOSTINGSLIST(dictionary, term(token))
10     ADDTOPOSTINGSLIST(postings_list, docID(token))
11 sorted_terms  $\leftarrow$  SORTTERMS(dictionary)
12 WRITEBLOCKTODISK(sorted_terms, dictionary, output_file)
13 return output_file
```

Merging of blocks is analogous to BSBI.



Single-pass in-memory indexing : compression

- 1 Compression makes SPIMI even more efficient.
 - Compression of terms
 - Compression of postings



Table of contents

- 1 Introduction
- 2 Sort-based index construction
- 3 Single-pass in-memory indexing (SPIMI)
- 4 Distributed indexing**
- 5 Dynamic indexing



Distributed indexing

- 1 For web-scale indexing: must use a distributed computer cluster
- 2 Individual machines are fault-prone.
Can unpredictably slow down or fail.
- 3 How do we exploit such a pool of machines?
- 4 Distributed index is partitioned across several machines - either according to term or according to document.




google data centers (Gartner estimates)

- 1 Google data centers mainly contain commodity machines. Data centers are distributed all over the world.
- 2 1 million servers, 3 million processors/cores
- 3 Google installs 100,000 servers each quarter.
- 4 Based on expenditures of 200250 million dollars per year. This would be 10% of the computing capacity of the world!
- 5 If in a non-fault-tolerant system with 1000 nodes, each node has 99.9% uptime, what is the uptime of the system (assuming it does not tolerate failures)?
- 6 Answer: 37%
- 7 Suppose a server will fail after 3 years. For an installation of 1 million servers, what is the interval between machine failures?
- 8 Answer: Less than two minutes.



Cluster architecture





Distributed indexing

- 1 Maintain a **master machine** directing the indexing job – considered "safe"
- 2 Break up indexing into sets of parallel tasks
- 3 Master machine assigns each task to an idle machine from a pool.



Parallel tasks

- 1 We will define two sets of parallel tasks and deploy two types of machines to solve them:
Parsers and **Inverters**
- 2 Break the input document collection into splits (corresponding to blocks in BSBI/SPIMI)
- 3 Each split is a subset of documents.



Parsers

- 1 Master assigns a split to an idle parser machine.
- 2 Parser reads a document at a time and emits (term,docID)-pairs.
- 3 Parser writes pairs into j term-partitions. Each for a range of terms first letters
E.g., a-f, g-p, q-z (here: $j = 3$)

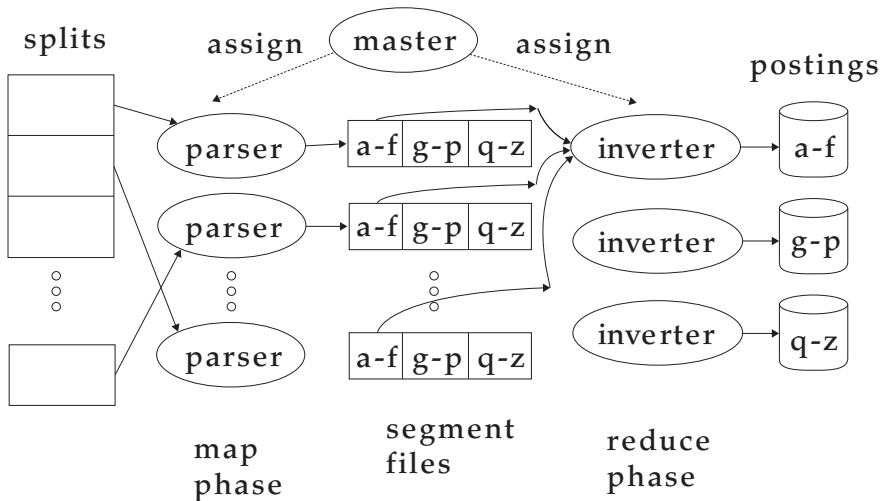


Inverters

- 1 An inverter collects all (term,docID) pairs (= postings) for one term-partition (e.g., for a-f).
- 2 Sorts and writes to postings lists



Data flow





Mapreduce

- 1 The index construction algorithm we just described is an instance of **MapReduce**.
- 2 MapReduce is a robust and conceptually simple framework for distributed computing . . . without having to write code for the distribution part.
- 3 The Google indexing system consisted of a number of phases, each implemented in MapReduce.
- 4 Index construction was just one phase.



MapReduce: word count example

```
map(key, value):  
  // key: document name; value: text of document  
  for each word w in value:  
    emit(w, 1)  
  
reduce(key, values):  
  // key: a word; value: an iterator over counts  
  result = 0  
  for each count v in values:  
    result += v  
  emit(result)
```



Table of contents

- 1 Introduction
- 2 Sort-based index construction
- 3 Single-pass in-memory indexing (SPIMI)
- 4 Distributed indexing
- 5 Dynamic indexing**



Dynamic indexing

- 1 Up to now, we have assumed that collections are **static**.
- 2 They rarely are: Documents are inserted, deleted and modified.
- 3 This means that the dictionary and postings lists have to be **dynamically** modified.



Dynamic indexing: simplest approach

- 1 Maintain **big main index on disk**
- 2 New docs go into **small auxiliary index in memory**.
- 3 Search across both, merge results
- 4 Periodically, merge auxiliary index into big index
- 5 Deletions:
 - Invalidation bit-vector for deleted docs
 - Filter docs returned by index using this bit-vector



Issues with auxiliary and main index

- 1 Frequent merges
- 2 Poor search performance during index merge



Logarithmic merge

- 1 Logarithmic merging amortizes the cost of merging indexes over time. Users see smaller effect on response times.
- 2 Maintain a series of indexes, each twice as large as the previous one.
- 3 Keep smallest (Z_0) in memory
- 4 Larger ones (I_0, I_1, \dots) on disk
- 5 If Z_0 gets too big ($\geq n$), write to disk as I_0
- 6 . . . or merge with I_0 (if I_0 already exists) and write merger to I_1 etc.



Logarithmic merge

LMERGEADDTOKEN(*indexes*, Z_0 , *token*)

```

1   $Z_0 \leftarrow \text{MERGE}(Z_0, \{token\})$ 
2  if  $|Z_0| = n$ 
3      then for  $i \leftarrow 0$  to  $\infty$ 
4          do if  $l_i \in \text{indexes}$ 
5              then  $Z_{i+1} \leftarrow \text{MERGE}(l_i, Z_i)$ 
6                  ( $Z_{i+1}$  is a temporary index on disk.)
7                   $\text{indexes} \leftarrow \text{indexes} - \{l_i\}$ 
8              else  $l_i \leftarrow Z_i$     ( $Z_i$  becomes the permanent index  $l_i$ .)
9                   $\text{indexes} \leftarrow \text{indexes} \cup \{l_i\}$ 
10                 BREAK
11      $Z_0 \leftarrow \emptyset$ 

```

LOGARITHMICMERGE()

```

1   $Z_0 \leftarrow \emptyset$     ( $Z_0$  is the in-memory index.)
2   $\text{indexes} \leftarrow \emptyset$ 
3  while true
4      do LMERGEADDTOKEN(indexes,  $Z_0$ , GETNEXTTOKEN())

```

Reading



Please read chapter 4 of Information Retrieval Book.