

Information Retrieval

Scores in a complete search system

Hamid Beigy

Sharif university of technology

October 19, 2018



Table of contents

- 1 Introduction
- 2 Improving scoring and ranking
- 3 A complete search engine



Introduction

- 1 We define **term frequency weight** of term t in document d as

$$tf_{t,d} = \sum_{x \in d} f_t(x) \text{ where } f_t(x) = \begin{cases} 1 & \text{if } x = t \\ 0 & \text{otherwise} \end{cases}$$

- 2 The log frequency weight of term t in d is defined as follows

$$w_{t,d} = \begin{cases} 1 + \log_{10} tf_{t,d} & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$$

- 3 We define the **idf weight** of term t as follows:

$$idf_t = \log_{10} \frac{N}{df_t}$$

- 4 We define the **tf-idf weight** of term t as **product of its tf and idf weights**.

$$w_{t,d} = (1 + \log tf_{t,d}) \cdot \log \frac{N}{df_t}$$



Cosine similarity between query and document

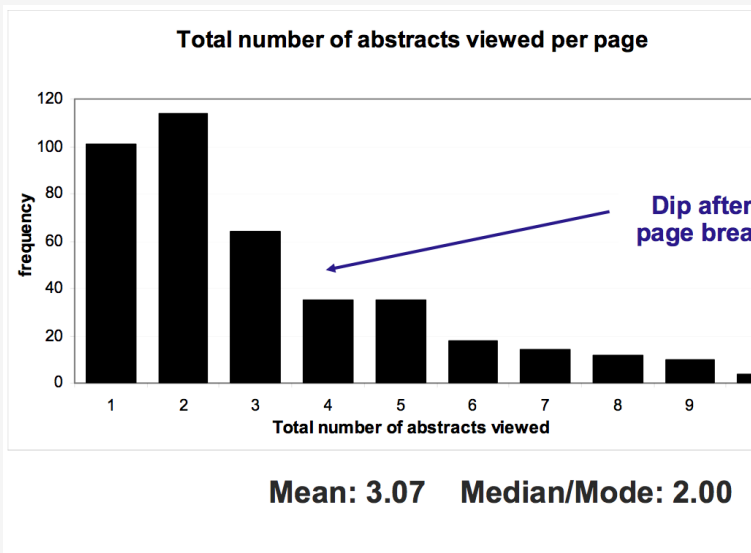
- 1 Cosine similarity between query q and document d is defined as

$$\cos(\vec{q}, \vec{d}) = \text{SIM}(\vec{q}, \vec{d}) = \frac{\vec{q}}{|\vec{q}|} \cdot \frac{\vec{d}}{|\vec{d}|} = \sum_{i=1}^{|V|} \frac{q_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2}} \cdot \frac{d_i}{\sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

- 2 q_i is the tf-idf weight of term i in the query.
- 3 d_i is the tf-idf weight of term i in the document.
- 4 $|\vec{q}|$ and $|\vec{d}|$ are the lengths of \vec{q} and \vec{d} .
- 5 $\vec{q}/|\vec{q}|$ and $\vec{d}/|\vec{d}|$ are length-1 vectors (= normalized).
- 6 Computing the cosine similarity is time-consuming task.




How many links do users view?





Looking versus clicking

- 1 Users view results two more often/ thoroughly.
- 2 Users click most frequently on result one.





Importance of ranking

- 1 **Viewing abstracts:** Users are a lot more likely to read the abstracts of the **top-ranked pages** (1, 2, 3, 4) than the abstracts of the **lower ranked pages** (7, 8, 9, 10).
- 2 **Clicking:** Distribution is even more skewed for clicking
- 3 In 1 out of 2 cases, users click on the **top-ranked page**.
- 4 Even if the top-ranked page is not relevant, 30% of users will click on it.
 - Getting the ranking right is very important.
 - Getting the top-ranked page right is most important



Table of contents

- 1 Introduction
- 2 Improving scoring and ranking**
- 3 A complete search engine



Speeding up document scoring

- 1 The scoring algorithm can be time consuming
- 2 Using heuristics can help saving time
- 3 **Exact top-score** vs **approximative top-score** retrieval

We can lower the cost of scoring by searching for K documents that are likely to be among the top-scores

- 4 General optimization scheme:
 - 1 find a set of documents A such that $K < |A| \ll N$, and whose is likely to contain many documents close to the top-scores
 - 2 return the K top-scoring document included in A



Index elimination

- 1 While processing the query, only consider terms whose idf_t exceeds a predefined threshold
Thus we avoid traversing the posting lists of high idf_t terms, lists which are generally long
- 2 Only consider documents where all query terms appear



Champion lists

- 1 We know which documents are the most relevant for a given term
- 2 For each term t , we pre-compute the list of the r most relevant (with respect to $w(t, d)$) documents in the collection
- 3 Given a query q , we compute

$$A = \bigcup_{t \in q} r(t)$$

r can depends on the document frequency of the term.



Static quality score

- 1 only consider documents which are considered as high-quality documents
- 2 Given a measure of quality $g(d)$, the posting lists are ordered by decreasing value of $g(d)$
- 3 Can be combined with champion lists, *i.e.* build the list of r most relevant documents wrt $g(d)$
- 4 Quality can be computed from the logs of users' queries



Impact ordering

- 1 Some sublists of the posting lists are of no interest
- 2 To reduce the time complexity:
 - query terms are processed by decreasing idf_t
 - postings are sorted by decreasing term frequency $tf_{t,d}$
 - Once idf_t gets low, we can consider only few postings
 - Once $tf_{t,d}$ gets smaller than a predefined threshold, the remaining postings in the list are skipped




Cluster pruning

- 1 The document vectors are gathered by proximity
- 2 We pick \sqrt{N} documents randomly \Rightarrow leaders
- 3 For each non-leader, we compute its nearest leader \Rightarrow followers
- 4 At query time, we only compute similarities between the query and the leaders
- 5 The set A is the closest document cluster
- 6 The document clustering should reflect the distribution of the vector space



Cluster pruning





Tiered indexes

- 1 This technique can be seen as a generalization of champion lists
- 2 Instead of considering one champion list, we manage layers of champion lists, ordered in increasing size:

index 1	l most relevant documents
index 2	next m most relevant documents
index 3	next n most relevant documents

Indexed defined according to thresholds

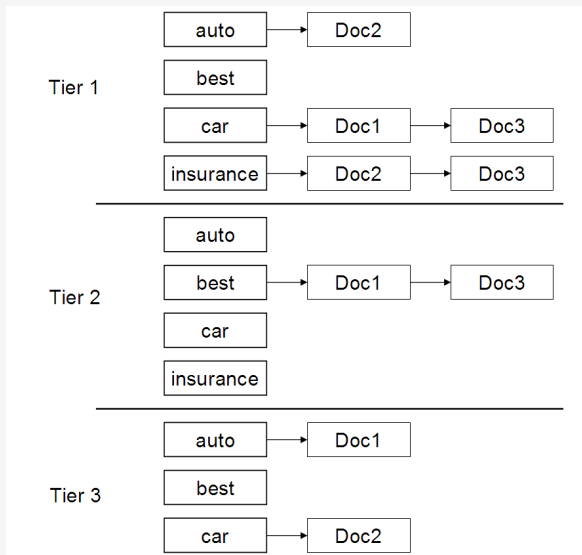
Dictionary	$v(\vec{d}_1)$	$v(\vec{d}_2)$	$v(\vec{d}_3)$
affection	0.996	0.993	0.847
jealous	0.087	0.120	0.466
gossip	0.017	0	0.254

$$\text{sim}(d_1, d_2) = 0.999$$

$$\text{sim}(d_1, d_3) = 0.888$$



Tiered indexes





Query-term proximity

- 1 Priority is given to documents containing many query terms in a close window
- 2 Needs to pre-compute n-grams
- 3 And to define a proximity weighting that depends on the window size n (either by hand or using learning algorithms)



Scoring optimizations – summary

- 1 Index elimination
- 2 Champion lists
- 3 Static quality score
- 4 Impact ordering
- 5 Cluster pruning
- 6 Tiered indexes
- 7 Query-term proximity



Table of contents

- 1 Introduction
- 2 Improving scoring and ranking
- 3 A complete search engine

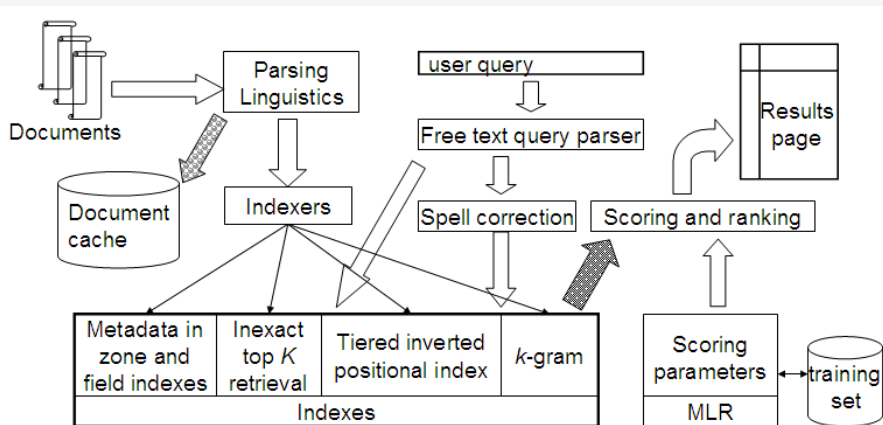


Putting it all together

- 1 Many techniques to retrieve documents (using logical operators, proximity operators, or scoring functions)
- 2 Adapted technique can be selected dynamically, by parsing the query
- 3 First process the query as a phrase query, if fewer than K results, then translate the query into phrase queries on bi-grams, if there are still too few results, finally process each term independently (real free text query)



A complete search engine



Reading



Please read chapter 7 of Information Retrieval Book.